


GPU-accelerated fractal imaging

Jeremy Ehrhardt


CS 81 - Spring 2009

Motivation

- Typical reconstruction filters blur edges, detail



Original



Reduced 4x, then enlarged 4x
with Mitchell filter

Motivation

- Fractal encoding
 - A **scale-independent** model of image
 - Find self-similar regions of image
 - Contractive transformations

Motivation

- Decoding fractal representation
 - Any scale
 - Larger than original image
 - Preserves appearance of high-frequency detail
 - Improved visual quality

Motivation

- Basic encoding algorithm
 - Divide image into small **range** blocks
 - Generate **domain** blocks from image somehow
 - For each range block
 - For each domain block
 - Find scale and offset for pixel luminance
 - $R = t(D) = s * D + o$
 - Pick domain with lowest mean squared error

Using the GPU

- Encoding highly parallelizable
 - Range selections don't depend on other ranges
 - Process multiple ranges at a time
 - Process multiple domains for a range at a time
 - Process multiple pixels in range, domain, or product at a time

Using the GPU

- Decoding also parallelizable
 - Short number of iterations (~10)
 - Take transform list from encoding
 - Decide on output size
 - For each pixel in output
 - Source, scale, offset don't change across iterations
 - Value depends only on image at previous iteration
 - Maps naturally to pixel shader

What I Did

- Software encoder and decoder
 - Python, using NumPy fast array extension
- GPU-assisted encoder
 - Most work done in GLSL pixel shaders
- Implementations similar
 - Easy debugging
 - All images were decoded in software

Encoding times

Encoder	Domain size	Range size	Image dimensions	Time (seconds)
Software	8	4	256 x 256	350
OpenGL	8	4	256 x 256	52
Software	4	2	256 x 256	5620
OpenGL	4	2	256 x 256	176

Domain and range sizes are pixels along edge of a square

- **CPU:** Intel Core 2 Duo @ 2.16 GHz
 - 2 GB RAM @ 0.667 GHz
- **GPU:** nVidia GeForce 7600 GT @ 0.56 GHz
 - 256 MB RAM @ 0.7 GHz

Results (no enlargement)

- Tradeoff between speckles and blocking




Fractal encoder (GPU, **4:2**)




Fractal encoder (GPU, **8:4**)

Results (no enlargement)




Fractal encoder (GPU, **4:2**)




Fractal encoder (GPU, **8:4**)

Results (no enlargement)



Fractal encoder (GPU, **4:2**)



Fractal encoder (GPU, **8:4**)

Results (2x enlargement)



Reduced 2x, then enlarged 2x
with fractal scaler (GPU, **4:2**)



Reduced 2x, then enlarged 2x
with fractal scaler (GPU, **8:4**)

Results (4x enlargement)



Reduced 4x, then enlarged 4x
with fractal scaler (GPU, **4:2**)



Reduced 4x, then enlarged 4x
with fractal scaler (GPU, **8:4**)

Side by Side



Reduced 2x, then enlarged 2x
with Mitchell filtering




Reduced 2x, then enlarged 2x
with fractal scaler (GPU, 8:4)

GPU vs CPU

- GPU has speckle noise
 - Probably due to lower accuracy



CPU, 8:4



GPU, 8:4

GPU vs CPU

- GPU has speckle noise
 - Probably due to lower accuracy



CPU, 4:2



GPU, 4:2

Accuracy issues

- Iterative nature of decoding process would tend to magnify errors
 - Testing on double-precision GPU needed

Wish list

- GPU-accelerated decoder
- CUDA or OpenCL port

Conclusions

- CPU encoder quality getting higher
 - Extremely slow
- GPU encoder is fast but has accuracy issues
 - Where are they?
 - Can they be fixed or compensated for?
 - CUDA or OpenCL port would help here